1 Requirement 1 : Correlation Detector

Simple correlation can be performed using a correlation of the template and image, all tests in this report used the obama.png as the template image. The results of using a simple correlation can be seen in figure 1. As can be seen this results is significant numbers of false positive detections, expecially in the flag section above the rows of people. Also many detections only cover half of a face.


Figure 1: Results of using Simple Correlation on g20 image

To improve this a normalised correlation technique was implemented. The normalised correlation is defined in the coursework specification and is shown in figure 4 for reference. This can be simplified in a series of steps to make use of correlations in order to improve performance. Firstly the value of $x - m_x 1$ can be precomputed as this is simple the template minus its mean. Next the numerator of the fraction can be simplified to $xy + m_x 1 m_y 1 - x m_y 1 - y m_x 1$, each of these four terms can then be computed using a correlation. These four correlations are shown in figure 3 where tempalate and image are the vectorised image and template and templateMean is the mean of the template and imageMean is a matrix of the mean values at each tempalte location, computed using another correlation with a template of values 1/1024 where the template size is 32 x 32.

$$r = \frac{(x - m_x 1)^T (y - m_y 1)}{|x - m_x 1||y - m_x 1|}$$

Figure 2: Normalised Correlation

```
1  MxY = filter2 (templateMean , image ) ;
2  XMy = filter2 (template , imageMean ) ;
3  MxMy = filter2 (templateMean , imageMean ) ;
4  XY = filter2 (template , image ) ;
```

Figure 3: Numerator Compute Code

Now the numerator of the normalised correlation is computed we can turn to the denominator. The denominator of the normalised correlation equation is the standard deviation of the tempalte and part of the image under the template [2]. The standard deviation of the template can be computed simply. For the part of the image under the template correlations can be used. The standard deviation can be rearranged to the equation shown in figure 4 where n is the number of pixels under the template, q is the squared sum of the pixels under the template and s is the sum of the pixels under the tempalte [1]. This can be computed using another three correlations to compute n, q and s as shown in figure 5 where the template is of size 32 x 32.

$$r = \frac{1}{n-1}(q - \frac{s^2}{n})$$

Figure 4: Standard Deviation

```
1  n = filter2(ones(size(32, 32), ones(32, 32), 'same');
2
3  imSquared = im .^2;
4  q = filter2(ones(32, 32), imSquared, 'same');
5
6  s = filter2(32, 32, image, 'same');
7
8  imStdDev = (q - ((s .^ 2) ./ n) ./ (n - 1));
9
10 imStdDev = sqrt(imStdDev);
```

Figure 5: Standard Deviation matlab

Using the normalised correlation the output of the face detection using obama.png as the template is shown in figure 6. As can be seen there is still a high number of false detections in the background of the image as well as in the tie area of some people. The next section will try and address these issues with an improved detector.

Figure 6: Results of using Normalised Correlation on g20 image

2 Requirement 2 : Improved Detector

The original normalised correlation detector achieves 27% overlap with the Viola-Jones face detection algorithm and this will be used as the performance measures for the improvements demonstrated in this section.

2.1 Guassian Blur

One improvement which was tried was to add a Gaussian blur using a correlation of a Gaussian template. The idea here was to remove some of the fine grain detain and high frequencies in the template image as these high frequencies may be causing some of the missed faces. This did increase the overlap to 29% the result can be seen in figure 7. The number of correctly detected faces is improved however there are still high numbers of false positives.


Figure 7: Results of using Normalised Correlation and Guassian Blur on g20 image

2.2 Colour Filtering

As many of the false detections are caused in the background of the image where the colours are not similar to skin tones a colour filter was implemented to try and reduce their occurrence. Using the training data supplied the upper and lower bounds for skin tone colours in red, green and blue where found. Any pixel which did not fall within these bounds was set to black. It was found that only

applying the colour filter to the g20 image and not the template caused highest nubmers of detections. This achived an overlap of 50%, which is shown in figure 8. As can be seen in the figure there are significantly less false detections in the background of the image and also more correct detections as a result of reducing the search space down to the pixels which are within skin tone ranges.


Figure 8: Results of using Normalised Correlation and Colour Filter on g20 image

2.3 Edge Detection

The final improvement on the original normalised correlation technique was using edge detection. The idea behind this was to simplfy the image down to just an edge detected image this means that the detection is just using the shape of the edges in the image as opposed to the colour or grayscale pixel value. It was expected that this would reduce the number of false detection is the background as the edges here are simple verticle lines which are not similar to those in the tempalate image. It was also expected that the number of correct detections would increase as the grey scale values are no longer being compared, so skin tone should not be an issue.

In order to create an edge detected image the Canny edge detection technique was applied. Firstly the Sobel Kernels are applied in two convolutions, for the horizontal and vertical. After this the gradients at each pixel are computed and are set to the closest of the two diagonals, horizontal and vertical. Finally the values are all thresholded to only leave the hard edges.

Using the edge detected image an overlap of 34% was achieved, this is shown in figure 9. As expected the number of false detections in the background of the image is reduced compared to the original normalised correlation and there are more positive detections.

Figure 9: Results of using Normalised Correlation and Edge Detection on g20 image

2.4 Harr Features

The final detection technique which was tried was to move away from using a convolution of a template image and to use Harr features like the Viola-Jones algorithm does **??**. The Viola-Jones algorithm uses a cascade of simple classifiers using a Harr feature to detect faces. Therefore the use of a set Harr features to detect an face should yeild an improved overlap with the Viola-Jones algorithm.

In order to do this a seven Harr Features were used each of a fixed size, these were chosen by running all possible sizes of the two, three and four segment Harr features over a test set of faces and non faces. The features which gave the highest success rate when run on the test set were used in the final implementation. An integral image of the target image is created in order to compare the Harr features too. Each of the Harr features were used as the template for a convolution across the Wintegral image. The average of each pixel in each convolution is calculated giving an matrix where high values represent where more of the Harr features matched the image below. An example of the output when run on the g20 image is shown in figure 10, as can be seen lighter areas are where faces are in the image. This technique yeilds a 49% overlap with the Viola-Jones algorithm.

Figure 10: Results of using Harr feature technique on g20 image

3   Requirement 3 : Nearest-Neighbour Classifier

In order to create a nearest-neighbour style classifier the euclidean distance between each of the detected face and the faces in the database are computed. Firstly each image in the database and the extracted images has its mean subtracted and is divided by its standard deviation in order to normalise lighting conditions. The euclidean distance is simply the squared distance between each pixel value in two images summed together. For each extracted image the distance between each image in the database is computed and the database image with the smallest distance is recorded and the class that image belongs to is the person that extraction is classified as. The code for this is in the EuclideanClassifier.m file.

Using this classifier 5 out of the 18 faces which are extracted by the face detection algorithm are correctly classified. This gives an overall success rate of 28%. This is shown in figure 11.

Figure 11: Results of face classification using Nearest-Neighbour Classifier

4   Requirement 4 : Improved Classifier

4.1   Using SVM and Naive Bayes classifiers

The first improvement which was tried was to replace the Nearest-Neighbour classifier with a different classifier. Two classifiers were tried first a Support Vector Machine and then a Naive Bayes classifier. Feature vectors for both classifiers are simply vectorised versions of the 32 x 32 images. The SVM approach comprises of a binary SVM for each of the classes of image. Each singe classifier returns true or false depending on whether the image passed is classified as the class the SVM was trained of. Each SVM is trained using all images of the face it is to classify as class True and all other faces in the database as class False. Then in order to classify am unknown face each SVM is run in turn until one returns true. The Naive Bayes approach uses a multi class Naive Bayes trained on the database of images provided.

Both of these classifiers produced very similar success rates of 28% for the SVM and 22% for the Naive Bayes these are no better than the Nearest-

Neighbour classifier. It was decided to change the features used in the classification next as opposed to modifying any of the classifiers parameters.

4.2   Using Eigen Faces

Eigen faces can be used as a feature for a face classifier. In order to create a set of eigen faces from the database firstly the average image is computed, the left most image in figure 12. The eigen faces are then computed by eigen value decomposition, each of the training images is placed in a M x N matrix where M is the number of images and N there vectorised length. The convolution matrix of this is then computed giving a N x N matrix, the eigen vectors of this are the eigen faces 12. Each eigen face is normalised between 0 and 1. The first five eigen faces are shown in figure 12.

Figure 12: Average Face and largest Eigen Faces, Left to Right

After the eigen faces have been computed then can be used to build feature vectors for images. This is done by computing the weights required for each of the eigen faces in order to reproduce the image. The weights are created by firstly taking the average face away from the target image then computing the dot product between each eigen face giving the weight for that eigen face. This gives a vector who's length is the same as the number of eigen faces, and can be used as a feature vector for face classification.

Using a set of training and testing data from the database it was found that the optimum eigen faces to use were the 50 largest. It was also found that a Naive Bayes classifier yielded the highest success rate. Using a naive bayes classifier trained with the weights for each of the images in the database and then run against the faces extracted by the Viola-Jones algorithm a success rate of 56significantly better than the nearest neighbor approach. The results are shown in figure 13.

Figure 13: Results of face classification using Eigen Faces

References

[1] MATLAB TRICKS. Calculate standard deviation using minimal memory. `http://matlabtricks.com/post-19/calculating-standard-deviation-using-minimal-memory`, 2013.

[2] WIKIPEDIA. Normalized cross-correlation. `https://en.wikipedia.org/wiki/Cross-correlation#Normalized_cross-correlation`, 2017.